Yelp and Crime

Kenneth Lin, Sid Naik, Tom McCormick

December 14, 2014

1 Problem Statement and Background

For our CS 194 final project, we decided to investigate the potential relationship between the City of San Francisco public safety data set and the data set provided by the Yelp API. With the recent civil unrest both inside and outside of the United States, and more recently, right here in Berkeley, we thought that it would be interesting to look into the factors that promote crime. One of our team members (Kenneth Lin) had also been robbed recently, so the problem is one that is dear to our hearts. Perhaps the most well-known correlation with crime rate is the income level of a neighborhood – the lower the income level, the higher the crime rate [1, p.93-94]. However, we wanted to show something more interesting. In particular, Yelp restaurants, in our experience, often reflect the wealth and well-being of its surrounding neighborhood – the presence of many highly rated restaurants, we believed, reflect the optimism in the economy of a neighborhood, as well as the wealth and "goodness" of that neighborhood. Therefore, we had conjectured that crime would negatively impact restaurant ratings, or that low restaurant ratings would be correlated with areas of high crime. We worked to show this throughout our project.

In particular, we wanted to know

✓ the effect of crime on restaurant ratings (or vice versa)

✓ the distribution of crime vs. the distribution of ratings

We had also wanted to predict crime density / severity using restaurant ratings or vice versa, but along the way we ran into issues of determining or implying causation in any of the methods we used. By using one to predict the other and trying to draw useful conclusions from this, we run the risk of assuming causation without definitive proof. There are many other problems that may arise as a result of this, which will be discussed in the **Lessons Learned** section.

1.1 City of San Francisco Public Safety Data Set

The City of San Francisco public safety data set is a record, written by the San Francisco Police Department, of incoming incident reports, either via phone

call, in person, or otherwise. The incidents are reported via the SFPD CABLE crime incident reporting system. The incidents are reported over the course of more than 11 years, from 1/1/2003 to present. Incidents are recorded when a police report is filled out during or after a crime incident. Crimes range from aggravated assault to vandalism to death reports.

A sample of records in the data set looks like the following:

	IncidntNum	Category	Descript	DayOfWeek	Date
0	140001966	NON-CRIMINAL	DEATH REPORT, CAUSE UNKNOWN	Wednesday	01/01/2014
1	140003025	NON-CRIMINAL	DEATH REPORT, CAUSE UNKNOWN	Thursday	01/02/2014
2	140004487	NON-CRIMINAL	DEATH REPORT, CAUSE UNKNOWN	Thursday	01/02/2014
3	140000059	NON-CRIMINAL	AIDED CASE	Wednesday	01/01/2014
4	140000071	VANDALISM	MALICIOUS MISCHIEF, VANDALISM OF VEHICLES	Wednesday	01/01/2014

Time	PdDistrict	Resolution	Location	X	Y
16:21	TENDERLOIN	NONE	400.0 Block of ELLIS ST	-122.413794	37.784772
02:00	MISSION	NONE	500.0 Block of JERSEY ST	-122.438235	37.750203
14:30	BAYVIEW	NONE	100.0 Block of CORAL CT	-122.371925	37.727898
00:17	SOUTHERN	NONE	1500.0 Block of MISSION ST	-122.417566	37.773892
00:30	SOUTHERN	NONE	0.0 Block of MARKET ST	-122.393966	37.795028

Figure 1: Sample of San Francisco crime data set

The majority of the fields are self-explanatory. However, there are a few things to note:

1. Category and descript are both categories, but category is more general. There are only 36 different "Categories" while there are 499 different "Descript"s in the year of 2014.

2. Resolution, though none are shown in the sample above, denote whether any action was taken and what that action was.

3. X denotes longitude, while Y denote latitude.

A more detailed analysis is in the attached `analysis.ipynb`.

1.2 Yelp Data Set

Yelp.com is a platform which publishes crowd-sourced reviews about local businesses. On Yelp, customers who have used the services of local businesses may write reviews of these businesses and provide ratings of their satisfaction. Reviewers may select from between 1 to 5 stars for each review they make, and a business's average rating is the average of the ratings of each of the reviews it has received. Yelp supplies a platform for all kinds of local businesses ranging from restaurants to barbers to museums; however, for the purpose of our research, we will look primarily at restaurants as they are a very large majority of the reviews on Yelp.

There are two primary ways to access the data on Yelp. First, we can utilize the search / business API (`http://www.yelp.com/developers/documentation`). The API provides a way to search for local businesses matching a particular key term ("restaurants", for example) near a geographical location, and get all the rating / review information about that restaurant. The API further allows us to narrow the search to only the geographically closest restaurants (not ranked by rating). This gives us a way to link the geographical location of crime incidents to the types of restaurants near that incident.

The other way of accessing Yelp data is through the academic data set (`https://www.yelp.com/academic_dataset`). The Yelp academic data set provides all the data and associated reviews of the 250 closest businesses to each of 30 universities, including UC Berkeley. Although not a random sample of all businesses on Yelp, the academic data set provides a much better estimate of all businesses in the Yelp data set population. For the purposes of the analysis, we will assume for now that this data set is a perfect sample of the entire Yelp data set, and that its average is indicative of the Yelp-wide average. Issues with this assumption will be addressed in the **Lessons Learned** section.

– basic analysis

2   Methods

2.1   Data Fetching

The bulk of our work was done in trying to get data from Yelp. As mentioned, there were two main ways that Yelp provides to access data, and those are the API and the academic data set.

To access the Yelp API, Yelp provided sample Python code (`https://github.com/Yelp/yelp-api/tree/master/v2/python`). However, as we found out, the sample code was buggy – not only did the code fail on certain calls, it even failed on the default call when the search term was "dinner" and the location was "San Francisco, CA". We contacted Yelp API support about this issue, but it seemed that the API wasn't well-maintained, and Yelp engineers didn't have time to update or fix the sample code. Therefore, we decided to fix the bug ourselves.

After much investigation, we figured out that whereas Python's `urllib` library encoded spaces into "+" characters, Yelp's server-side authentication expected the OAuth-signed URLs to use "%20" as the proper encoding for space. In this context (the query arguments in a URL), both should be valid, but Yelp's authenticator only expected the latter. After figuring this out, we let the Yelp engineers know of the bug, and were able to begin building a temporary work-around to fetch our data.

We also encountered other problems in data fetching. In addition to the bugs in the API, there were rate limiting and quantity limiting issues as well. In particular, we could only access 20 results at a time, and a maximum of 40 total results for searches that sorted the results by only distance or rating. This limited the ways in which we could obtain and analyze Yelp data. Further, we weren't able to get as much data as we needed (because of Yelp's rate limiting). However, it was still possible to search for the closest restaurants near any location specified by longitude and latitude. As a result, we decided to search for the 20 closest restaurants near each crime and look at the characteristics of this set of restaurants.

In the end, we were able to build a fault-tolerant work-around to Yelp's broken authentication system. We queried Yelp's API for the 20 closest restaurants near each crime in the City of San Francisco crime data set for a total of 12,000 crimes (9704 unique incidents due to multiple criminal offenses per incident). These results were stored in our MongoDB instance that we installed on our EC2 server.

2.2 Visualization

2.3 Looking for Correlations

Our primary purpose (initially) was to determine how crime affected restaurant ratings. Specifically, given the limitations of the data we could get from Yelp, we wanted to see whether Yelp restaurant ratings were any different when they were near crime hotspots.

At this point, we could access two data sets – the data set of "near-crime" restaurants, obtained through querying for restaurants near crime incident locations via the Yelp API, and the Yelp academic data set. Given this data, we had a few options to draw a correlation:

1. We could look at the "average rating" of a crime – that is, we could take the average rating of the 20 closest restaurants to a crime incident and call that the average rating of restaurants near that crime. Then, we could compare that rating with the Yelp-wide average from the academic data set.

2. We could look at all the *restaurants* near any of the crime incidents in the crime data set. Then, we can put these together as a set of "near-crime" restaurants, and look at the rating distribution of these compared to the academic data set.

We had originally planned to execute option 1. However, as we realized from exploring the crime data set, crime density is very different for different areas of San Francisco. If we look at the average rating of each crime, it is likely that the results would be very heavily weighted by the ratings of restaurants in high crime density areas (ie. Market Street). Although our very purpose was to show that high crime density areas were prone to different kinds of Yelp ratings, we did not want to artificially induce this by weighting the results by the crime density itself. Therefore, we chose the second option. Surprisingly, we discovered that in the 9704 unique incidents we looked at and the 194080 restaurants near those incidents, there were only a total of 658 unique restaurants. This implies that the crimes were indeed very clustered, and that use of the "average rating of a crime" method (option 1) would have resulted in significant bias to restaurants near the clustered areas. This will be discussed more in the **Lessons Learned** section.

To determine whether there were any significant differences between the two data sets, we used a Student's t-test as our statistical hypothesis test. In this case, the null hypothesis was that the two data sets, both of Yelp ratings in San Francisco, have identical expected values (means). The t-test results are more thoroughly discussed in the **Results** section. Through the analysis, however, we learned instead that restaurants near high-crime areas actually had higher ratings than the Yelp-wide average.

However, as we learned through

2.4   Accounting for Population Density

As a result of our findings, we realized that causation != correlation.
   – tract difficulties, calculating tract centers

2.5   Predicting Crimes

3   Tools

✓ pandas – Pandas was our primary means of manipulating data sets. We used Pandas in our data fetching, data cleaning, and data analysis process.

✓ MongoDB – We installed MongoDB on the EC2 server and used it primarily for storing Yelp data. MongoDB was a great choice because the document-based nature of MongoDB was perfect for storing all the JSON documents we received from the Yelp API. In addition, we weren't sure what kind of data we needed to store when we first began fetching data, so using Mongo allowed us to easily keep all of the data in case we needed more than what we had thought.

✓ pyMongo – pyMongo was our choice Python interface to MongoDB.

✓ numpy – NumPy was used in conjunction with both Pandas and SciPy.

✓ scipy – We used SciPy mostly for its large library of statistical methods (t-tests, etc.).

✓ D3.js – We used D3 in conjunction with other frameworks when visualizing the data.

✓ Google Maps API – The Google Maps API provided a base for many of the location-based visualizations we needed to create. In addition, the Google Maps API also had a heatmap plugin, which we tried in addition to heatmap.js.

✓ heatmap.js – We primarily used heatmap.js for our heatmap visualizations.

4 Results

A preliminary analysis of just the public safety data set is available in the attached `analysis.pynb`. Below follows our results and conclusions from looking at Yelp restaurant rating data near crime hotspots.

5 Putting it together

The results of the t-test between the near-crime data set and subsets of the academic data set are detailed in Table 1.

Data set	Mean	Variance	t-statistic*	p-value*
Near-crime	4.072	0.107		
All businesses (all)	3.618	0.871	-12.441	2.39×10^{-35}
All businesses (Berkeley)	3.629	0.701	-12.390	3.43×10^{-33}
All businesses (Stanford)	3.696	0.716	-8.310	4.34×10^{-16}
All businesses (Los Angeles)	3.571	0.976	-12.542	1.66×10^{-34}
Restaurants only (all)	3.482	0.545	-20.291	7.10×10^{-89}
Restaurants only (Berkeley)	3.413	0.382	-20.490	9.34×10^{-77}
Restaurants only (Stanford)	3.307	0.270	-14.367	3.30×10^{-41}
Restaurants only (Los Angeles)	3.431	0.598	-18.797	2.28×10^{-68}

Table 1: t-test results against "near-crime" data set

* against the "near-crime" data set
– whisker plots
– distribution – ratings
– CONCLUSIONS??!?
– map visualization problems —- yelp data biased to near crime —- not enough data on all of san francisco to create proper viz
– condition tests in certain neighborhoods

6 Lessons Learned

– Yelp API shit; no one uses – data cleaning hard – causation and correlation – tract difficulties – clear problem statement —-¿ not our fault, more because we didn't have experience with data science projects and no proper guidance – problems with the data (not sure where it's from) – don't draw conclusions before the results are out – note about Yelp data not having "price" characteristic – academic data set biased – keyword match problems

References

[1] S. D. Levitt, "The Changing Relationship between Income and Crime Victimization," September 1999. [Online]. Available: http://pricetheory. uchicago.edu/levitt/Papers/LevittTheChangingRelationship1999.pdf